

class07

Joshua Martin (PID: A18545389)


Today we will explore some fundamental machine learning methods including clustering and dimensionality reduction.

K-means clustering

To see how this works let's first makeup some data to cluster where we know what the answer should be. We can use the `rnorm()` function to help here:


```
hist( rnorm(500, mean=5))
```

Histogram of `rnorm(500, mean = 5)`



```
x <- c( rnorm(30, mean=-3), rnorm(30,mean=3))  
y <- rev(x)
```

```
x <- cbind(x,y)  
plot(x)
```



The function for K-means clustering in “base” R is `kmeans()`

```
k <- kmeans(x, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1 -3.205319  2.960707
2  2.960707 -3.205319

```

Clustering vector:

```
Within cluster sum of squares by cluster:  
[1] 53.30216 53.30216  
(between_SS / total_SS =  91.5 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

To get at the results of the returned list object we can use the dollar \$ syntax

Q. How many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What component of your result object details - cluster assignment/membership?
- cluster center?

```
head(k$cluster)
```


```
[1] 1 1 1 1 1 1
```

```
k$centers
```

| | x | y |
|---|-----------|-----------|
| 1 | -3.205319 | 2.960707 |
| 2 | 2.960707 | -3.205319 |

Q. Make a clustering results figure of the data colored by cluster membership.

```
plot(x, col=k$cluster, pch=16)  
points(k$centers, col="blue", pch=15, cex=2)
```



K-means clustering is very popular as it is very fast and relatively straight forward: it takes numeric data input and returns the cluster membership vector etc.

The “issue” is we tell `kmeans()` how many clusters we want!

Q. Run kmeans again and cluster into 4 grps/clusters and plot the results like we did above?

```
k4 <- kmeans(x, centers = 4)
k4
```

K-means clustering with 4 clusters of sizes 13, 17, 13, 17

Cluster means:

| | x | y |
|---|-----------|-----------|
| 1 | -3.743862 | 2.130260 |
| 2 | 3.595756 | -2.793493 |
| 3 | 2.130260 | -3.743862 |
| 4 | -2.793493 | 3.595756 |

Clustering vector:

```
[1] 4 1 4 4 4 4 4 1 1 1 1 1 1 4 4 1 1 4 1 4 4 1 4 4 1 4 1 4 2 3 2 3 2 2 3 2
[39] 2 3 2 3 3 2 3 2 2 3 3 3 3 2 2 2 2 2 3 2
```

```
Within cluster sum of squares by cluster:  
[1] 20.73095 10.09639 20.73095 10.09639  
(between_SS / total_SS =  95.1 %)
```


Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
k4$size
```

```
[1] 13 17 13 17
```

```
plot(x, col = k4$cluster, pch = 16)  
points(k4$centers, col="blue", pch = 15, cex = 2)
```



Q. Generate a scree plot.


```

# Basic scree plot
wss <- numeric(10)

for (k in 1:10) {
  wss[k] <- kmeans(x, centers = k)$tot.withinss
}

plot(1:10, wss, type = "b",
      xlab = "Number of clusters (k)",
      ylab = "Total within-cluster sum of squares",
      main = "Scree Plot")

```



Hierarchical Clustering

The main “base” R function for Hierarchical Clustering is called `hclust()`. Here we can’t just input our data we need to first calculate a distance matrix (e.g., `dist()`) for our data and use this as input to `hclust()`

```


d <- dist(x)
hc <- hclust(d)
hc

```

```
Call:  
hclust(d = d)  
  
Cluster method : complete  
Distance       : euclidean  
Number of objects: 60
```

There is a plot method for hclust results lets try it

```
plot(hc)
abline(h=10, col="red")
```




To get our cluster “membership” vector (i.e. our main clustering result) we can “cut” the tree at a given height or at a height that yields a given “ k ” groups.

```
cutree(hc, h=10)
```

```
grps <- cutree(hc, k=2)
```

Q. Plot the data with our hclust result coloring

```
plot(x, col=grps)
```



Principal Component Analysis (PCA)

PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
head(x)
```

| | X | England | Wales | Scotland | N.Ireland |
|---|--------------|---------|-------|----------|-----------|
| 1 | Cheese | 105 | 103 | 103 | 66 |
| 2 | Carcass_meat | 245 | 227 | 242 | 267 |
| 3 | Other_meat | 685 | 803 | 750 | 586 |
| 4 | Fish | 147 | 160 | 122 | 93 |

```

5 Fats_and_oils      193   235      184      209
6           Sugars     156   175      147      139

```

```

rownames(x) <- x[,1]
x <- x[,-1]
x

```

| | England | Wales | Scotland | N.Ireland |
|--------------------|---------|-------|----------|-----------|
| Cheese | 105 | 103 | 103 | 66 |
| Carcass_meat | 245 | 227 | 242 | 267 |
| Other_meat | 685 | 803 | 750 | 586 |
| Fish | 147 | 160 | 122 | 93 |
| Fats_and_oils | 193 | 235 | 184 | 209 |
| Sugars | 156 | 175 | 147 | 139 |
| Fresh_potatoes | 720 | 874 | 566 | 1033 |
| Fresh_Veg | 253 | 265 | 171 | 143 |
| Other_Veg | 488 | 570 | 418 | 355 |
| Processed_potatoes | 198 | 203 | 220 | 187 |
| Processed_Veg | 360 | 365 | 337 | 334 |
| Fresh_fruit | 1102 | 1137 | 957 | 674 |
| Cereals | 1472 | 1582 | 1462 | 1494 |
| Beverages | 57 | 73 | 53 | 47 |
| Soft_drinks | 1374 | 1256 | 1572 | 1506 |
| Alcoholic_drinks | 375 | 475 | 458 | 135 |
| Confectionery | 54 | 64 | 62 | 41 |

```
x <- read.csv("UK.foods.csv")
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
head(x)
```

| X | England | Wales | Scotland | N.Ireland |
|-----------------|---------|-------|----------|-----------|
| 1 Cheese | 105 | 103 | 103 | 66 |
| 2 Carcass_meat | 245 | 227 | 242 | 267 |
| 3 Other_meat | 685 | 803 | 750 | 586 |
| 4 Fish | 147 | 160 | 122 | 93 |
| 5 Fats_and_oils | 193 | 235 | 184 | 209 |
| 6 Sugars | 156 | 175 | 147 | 139 |

```
dim(x)
```

```
[1] 17 5
```

```
rownames(x) <- x[,1]  
x <- x[,-1]
```

```
dim(x)
```

```
[1] 17 4
```


```
head(x)
```

| | England | Wales | Scotland | N.Ireland |
|---------------|---------|-------|----------|-----------|
| Cheese | 105 | 103 | 103 | 66 |
| Carcass_meat | 245 | 227 | 242 | 267 |
| Other_meat | 685 | 803 | 750 | 586 |
| Fish | 147 | 160 | 122 | 93 |
| Fats_and_oils | 193 | 235 | 184 | 209 |
| Sugars | 156 | 175 | 147 | 139 |

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


I prefer using row.names=1 argument in read.csv() when importing the data. More robust, cleaner.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3. Changing what optional argument in the above barplot() function results in the following plot? beside=FALSE


```
barplot(as.matrix(x), beside = FALSE, col = rainbow(nrow(x)))
```



Q4. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

If a point lies on the diagonal, it means that food is consumed equally in both countries and their values are similar or the same.

```
pairs(x, col=rainbow(10), pch=16)
```



Main point: It can be difficult to spot major trends and patterns even in relatively small multivariate datasets (here we only have 17 dimensions, typically we have 1000s)

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`

I will take the transpose of our data so the “foods” are in the columns:

```
pca <- prcomp( t(x) )
summary(pca)
```


```
Importance of components:
```

| | PC1 | PC2 | PC3 | PC4 |
|------------------------|----------|----------|----------|-----------|
| Standard deviation | 324.1502 | 212.7478 | 73.87622 | 2.921e-14 |
| Proportion of Variance | 0.6744 | 0.2905 | 0.03503 | 0.000e+00 |
| Cumulative Proportion | 0.6744 | 0.9650 | 1.00000 | 1.000e+00 |

```
pca$x
```


| | PC1 | PC2 | PC3 | PC4 |
|-----------|------------|-------------|------------|---------------|
| England | -144.99315 | -2.532999 | 105.768945 | -9.152022e-15 |
| Wales | -240.52915 | -224.646925 | -56.475555 | 5.560040e-13 |
| Scotland | -91.86934 | 286.081786 | -44.415495 | -6.638419e-13 |
| N.Ireland | 477.39164 | -58.901862 | -4.877895 | 1.329771e-13 |

```
cols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=cols, pch=16)
```




```
library(ggplot2)
```

```
ggplot(pca$x) +  
  aes(PC1, PC2) +  
  geom_point(col=cols)
```



```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```



PCA looks super useful and will come back to describe this further next day :-)