

class12

Joshua Martin (PID: A18545389)

Table of contents

| | |
|--|----|
| Background | 1 |
| Data import | 1 |
| Toy differential gene expression | 3 |
| Volcano Plot | 10 |
| Save our results | 12 |

Background

Today we will analyze some RNASeq data from Himes et al. on the effects of a common steroid (dexamethasone) on airway smooth muscle cells (ASM cells).

Our starting point is the “counts” data and “metadata” that contain the count values for each gene in their different experiment (i.e. cell lines with or without the drug).

Data import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Let's have a wee peak at these objects:

```
head(counts)
```

| | SRR1039508 | SRR1039509 | SRR1039512 | SRR1039513 | SRR1039516 |
|------------------|------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 486 | 904 | 445 | 1170 |
| ENSG000000000005 | 0 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 523 | 616 | 371 | 582 |
| ENSG000000000457 | 347 | 258 | 364 | 237 | 318 |
| ENSG000000000460 | 96 | 81 | 73 | 66 | 118 |
| ENSG000000000938 | 0 | 0 | 1 | 0 | 2 |
| | SRR1039517 | SRR1039520 | SRR1039521 | | |
| ENSG000000000003 | 1097 | 806 | 604 | | |
| ENSG000000000005 | 0 | 0 | 0 | | |
| ENSG000000000419 | 781 | 417 | 509 | | |
| ENSG000000000457 | 447 | 330 | 324 | | |
| ENSG000000000460 | 94 | 102 | 74 | | |
| ENSG000000000938 | 0 | 0 | 0 | | |

Q. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

```
metadata
```

| | id | dex | celltype | geo_id |
|---|------------|---------|----------|------------|
| 1 | SRR1039508 | control | N61311 | GSM1275862 |
| 2 | SRR1039509 | treated | N61311 | GSM1275863 |
| 3 | SRR1039512 | control | N052611 | GSM1275866 |
| 4 | SRR1039513 | treated | N052611 | GSM1275867 |
| 5 | SRR1039516 | control | N080611 | GSM1275870 |
| 6 | SRR1039517 | treated | N080611 | GSM1275871 |
| 7 | SRR1039520 | control | N061011 | GSM1275874 |
| 8 | SRR1039521 | treated | N061011 | GSM1275875 |

```
ncol(counts)
```

```
[1] 8
```

```
nrow(counts)
```

```
[1] 38694
```

Q2. How many ‘control’ cell lines do we have?

```
metadata$dex == "control"  
  
[1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE  
  
sum(metadata$dex == "control")  
  
[1] 4
```

Toy differential gene expression

To start our analysis let’s calculate the mean counts for all genes in the “control” experiments.

1. Extract all “control” columns for the `counts` object
 2. Calculate the mean for all rows (i.e. genes) of these “control” columns
- 3-4. Do the same for “treated” 5. Compare these `control.mean` and `treated.mean` values.

```
control inds <- metadata$dex == "control"  
control counts <- counts[, control inds]  
head(control counts)
```

| | SRR1039508 | SRR1039512 | SRR1039516 | SRR1039520 |
|------------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 904 | 1170 | 806 |
| ENSG000000000005 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 616 | 582 | 417 |
| ENSG000000000457 | 347 | 364 | 318 | 330 |
| ENSG000000000460 | 96 | 73 | 118 | 102 |
| ENSG000000000938 | 0 | 1 | 2 | 0 |

```
control means <- rowSums( control counts )/4  
head(control means)
```

| ENSG000000000003 | ENSG000000000005 | ENSG000000000419 | ENSG000000000457 | ENSG000000000460 |
|------------------|------------------|------------------|------------------|------------------|
| 900.75 | 0.00 | 520.50 | 339.75 | 97.25 |
| ENSG000000000938 | | | | |
| 0.75 | | | | |

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

use rowMeans instead of dividing by 4 after setting by group size.

```
control.mean <- rowMeans(counts[, metadata$dex == "control", drop = FALSE])  
treated.mean <- rowMeans(counts[, metadata$dex == "treated", drop = FALSE])  
sd
```


```
treated inds <- metadata$dex == "treated"  
  
treated.counts <- counts[ , treated inds]  
  
treated.means <- rowMeans(treated.counts)
```

Store these together for ease of bookkeeping as `meancounts`

```
meancounts <- data.frame(control.means, treated.means)  
head(meancounts)
```

| | control.means | treated.means |
|------------------|---------------|---------------|
| ENSG000000000003 | 900.75 | 658.00 |
| ENSG000000000005 | 0.00 | 0.00 |
| ENSG000000000419 | 520.50 | 546.00 |
| ENSG000000000457 | 339.75 | 316.50 |
| ENSG000000000460 | 97.25 | 78.75 |
| ENSG000000000938 | 0.75 | 0.00 |

```
plot(meancounts[,1], meancounts[,2], xlab="Control", ylab="Treated")
```

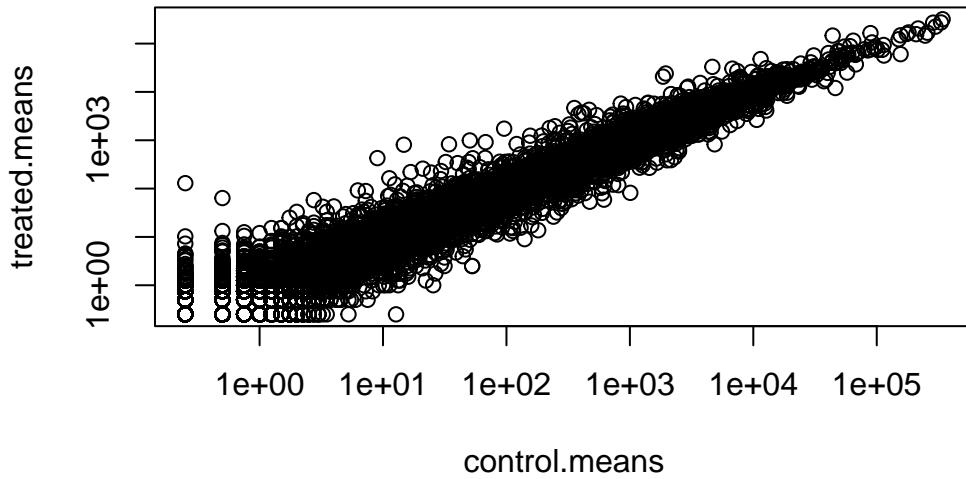


Make this a log log plot

```
plot(meancounts, log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We often talk metrics like “log2 fold-change”

```
# control/treated
log2(10/10)
```

```
[1] 0
```

```
log2(10/20)
```

```
[1] -1
```

```
log2(20/10)
```

```
[1] 1
```

```
log2(10/40)
```

```
[1] -2
```

```

zero.vals <- which(meancounts[,1:2] == 0, arr.ind = TRUE)
to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm, ]

```

What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

returns both the row and column positions of zeros. we take the row first column to remove genes with any zero counts, and unique() prevents removing the same row multiple times.

Let's calculate the log2 fold change for our treated over control mean counts.

```

meancounts$log2fc <-
log2(meancounts$treated.means /
    meancounts$control.means)

```

```
head(meancounts)
```

| | control.means | treated.means | log2fc |
|------------------|---------------|---------------|-------------|
| ENSG000000000003 | 900.75 | 658.00 | -0.45303916 |
| ENSG000000000005 | 0.00 | 0.00 | NaN |
| ENSG00000000419 | 520.50 | 546.00 | 0.06900279 |
| ENSG00000000457 | 339.75 | 316.50 | -0.10226805 |
| ENSG00000000460 | 97.25 | 78.75 | -0.30441833 |
| ENSG00000000938 | 0.75 | 0.00 | -Inf |

A common “rule of thumb” is a log2 fold change of +2 and -2 to call genes “Up regulated” or “Down regulated”.

```
sum(meancounts$log2fc > +2, na.rm=T)
```

```
[1] 1846
```

Number of “down” genes at -2 threshold

```
sum(meancounts$log2fc <= -2, na.rm=T)
```

```
[1] 2330
```

The above data is missing a statistical significance test.

```
up.ind <- mycounts$log2fc > 2
down.ind <- mycounts$log2fc < (-2) ## DESeq2 analysis
```

```
up.ind <- mycounts$log2fc > 2
down.ind <- mycounts$log2fc < (-2)
```

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

```
sum(mycounts$log2fc > 2, na.rm = TRUE)
```

```
[1] 0
```

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
sum(mycounts$log2fc < -2, na.rm = TRUE)
```

```
[1] 0
```

10. Do you trust these results? Why or why not?

No. they are calculated from unnormalized mean counts and not accounting for variance, may include false positives/negatives.

Let's do this analysis properly and keep our inner stats nerd happy = i.e. are the differences we see between drug and no drug statistically significant given the replicate experiments?

```
library(DESeq2)
```

For DESeq analysis we need three things

- count values (`countData`)
- metadata telling us about the columns in `countData` (`colData`)

Our first function from DESeq 2 will setup the input required for analysis by storing all these 3 things together.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                               colData = metadata,
                               design = ~dex)
```

```
converting counts to integer mode
```

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in  
design formula are characters, converting to factors
```

The main function in DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

```
estimating size factors
```

```
estimating dispersions
```

```
gene-wise dispersion estimates
```

```
mean-dispersion relationship
```

```
final dispersion estimates
```

```
fitting model and testing
```

```
results(dds)
```

```
log2 fold change (MLE): dex treated vs control
```

```
Wald test p-value: dex treated vs control
```

```
DataFrame with 38694 rows and 6 columns
```

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|-----------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.1942 | -0.350703 | 0.168242 | -2.084514 | 0.0371134 |
| ENSG000000000005 | 0.0000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.1342 | 0.206107 | 0.101042 | 2.039828 | 0.0413675 |
| ENSG000000000457 | 322.6648 | 0.024527 | 0.145134 | 0.168996 | 0.8658000 |
| ENSG000000000460 | 87.6826 | -0.147143 | 0.256995 | -0.572550 | 0.5669497 |
| ... | ... | ... | ... | ... | ... |
| ENSG00000283115 | 0.000000 | NA | NA | NA | NA |
| ENSG00000283116 | 0.000000 | NA | NA | NA | NA |
| ENSG00000283119 | 0.000000 | NA | NA | NA | NA |
| ENSG00000283120 | 0.974916 | -0.66825 | 1.69441 | -0.394385 | 0.693297 |
| ENSG00000283123 | 0.000000 | NA | NA | NA | NA |

```

      padj
<numeric>
ENSG000000000003 0.163017
ENSG000000000005     NA
ENSG000000000419 0.175937
ENSG000000000457 0.961682
ENSG000000000460 0.815805
...
ENSG0000283115     NA
ENSG0000283116     NA
ENSG0000283119     NA
ENSG0000283120     NA
ENSG0000283123     NA

```

```
36000 * 0.05
```

```
[1] 1800
```

Volcano Plot

```

res <- results(dds)
head(res)

```

```


log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 6 columns
  baseMean log2FoldChange      lfcSE      stat    pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000      NA       NA       NA       NA
ENSG000000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG000000000460 87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938 0.319167 -1.732289  3.493601 -0.495846 0.6200029
      padj
<numeric>
ENSG000000000003 0.163017
ENSG000000000005     NA
ENSG000000000419 0.175937
ENSG000000000457 0.961682

```


```
ENSG00000000460 0.815805
ENSG00000000938 NA
```

This is a common summary result figure from these types of experiments and plot the log2 fold-change vs the adjusted p-value

```
plot(res$log2FoldChange, res$padj)
```



```
plot( res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="red")
abline(h=-log(0.04), col="red")
```



Save our results

```
write.csv(res, file="my_results.csv")
```